

Create an "Academic performance" dataset of students and perform the following operations using Python.

1. Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them.
2. Scan all numeric variables for outliers. If there are outliers, use any of the suitable techniques to deal with them.
3. Apply data transformations on at least one of the variables. The purpose of this transformation should be one of the following reasons: to change the scale for better understanding of the variable, to convert a non-linear relation into a linear one, or to decrease the skewness and convert the distribution into a normal distribution. Reason and document your approach properly.

```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Read the data from the CSV file
data = pd.read_csv('data.csv')

print(data.head())
```

	Roll Number	First Name	Last Name	Mobile Number	CGPA1	CGPA2	age
0	261	Rishi	Gupta	919954629666	9.92	5.06	19
1	827	Seema	Singh	919085484267	0.22	2.83	20
2	566	Pooja	Rao	919692747629	4.00	5.78	20
3	431	Vikram	Trivedi	919289900918	3.93	NaN	20
4	688	Rishi	Nair	918510327681	0.34	3.35	22

```
In [ ]: # 1. Scan all variables for missing values and inconsistencies. If there are mis
# inconsistencies, use any of the suitable techniques to deal with them.

# Check for missing values
print(data.isnull().sum())
```

Roll Number	0
First Name	0
Last Name	0
Mobile Number	0
CGPA1	5
CGPA2	3
age	0

dtype: int64


```
In [ ]: # Check for inconsistencies
print(data.describe())
```

	Roll Number	Mobile Number	CGPA1	CGPA2	age
count	50.000000	5.000000e+01	45.000000	47.000000	50.000000
mean	492.440000	9.185083e+11	5.092444	4.586170	22.780000
std	250.705792	9.131771e+08	3.068496	2.688698	22.994489
min	15.000000	9.170212e+11	0.000000	0.250000	2.000000
25%	288.000000	9.176502e+11	2.920000	2.730000	18.000000
50%	527.500000	9.186161e+11	4.750000	4.080000	19.500000
75%	659.000000	9.192426e+11	8.110000	6.580000	22.750000
max	994.000000	9.199546e+11	9.920000	9.720000	180.000000

```
In [ ]: # fill missing values with the mean of the column on CGPA1 and CGPA2
data['CGPA1'] = data['CGPA1'].fillna(data['CGPA1'].mean())
data['CGPA2'] = data['CGPA2'].fillna(data['CGPA2'].mean())
```

```
In [ ]: # 2. Scan all numeric variables for outliers. If there are outliers, use any of
# to deal with them.

# Check for outliers
sns.boxplot(data['age'])
plt.show()
```



```
In [ ]: # look for outliers in the age column
Q1 = data['age'].quantile(0.25)
Q3 = data['age'].quantile(0.75)

IQR = Q3 - Q1

print("Q1: ", Q1)
print("Q3: ", Q3)
print("IQR: ", IQR)
```

```
Q1: 18.0
Q3: 22.75
IQR: 4.75
```

```
In [ ]: # print the number of outliers
outliers = data[(data['age'] < (Q1 - 1.5 * IQR)) | (data['age'] > (Q3 + 1.5 * IQR))]
print(outliers)
```

	Roll Number	First Name	Last Name	Mobile Number	CGPA1	CGPA2	age
6	532	Vaishnavi	Jha	919009969408	5.63	4.38	2
27	722	Vaishnavi	Patel	917550452611	9.49	3.29	180

```
In [ ]: # replace outliers with the mode

data['age'] = data['age'].mask(data['age'] > Q3 + 1.5 * IQR, data['age'].mode()[0])
```

```
data['age'] = data['age'].mask(data['age'] < Q1 - 1.5 * IQR, data['age'].mode()[0])
print(data['age'])
```

```
0    19
1    20
2    20
3    20
4    22
5    19
6    16
7    23
8    18
9    24
10   17
11   16
12   19
13   22
14   23
15   21
16   16
17   25
18   16
19   24
20   19
21   18
22   18
23   23
24   22
25   23
26   20
27   16
28   24
29   16
30   23
31   24
32   23
33   18
34   16
35   18
36   16
37   22
38   22
39   19
40   23
41   19
42   22
43   16
44   17
45   22
46   19
47   16
48   16
49   19
```

Name: age, dtype: int64

```
In [ ]: # 3. Apply data transformations on at least one of the variables. The purpose of
# should be one of the following reasons: to change the scale for better underst
# variable, to convert a non-linear relation into a linear one, or to decrease t
# convert the distribution into a normal distribution.
```

```
# Reason and document your approach properly.

# The age column has a centered data. We can apply a log transformation to the a
# convert the distribution into a normal distribution.

# log transformation
data['age'] = data['age'].apply(lambda x: np.log(x) if x > 0 else 0)

# display the transformed data
print(data['age'])
```

0	2.944439
1	2.995732
2	2.995732
3	2.995732
4	3.091042
5	2.944439
6	2.772589
7	3.135494
8	2.890372
9	3.178054
10	2.833213
11	2.772589
12	2.944439
13	3.091042
14	3.135494
15	3.044522
16	2.772589
17	3.218876
18	2.772589
19	3.178054
20	2.944439
21	2.890372
22	2.890372
23	3.135494
24	3.091042
25	3.135494
26	2.995732
27	2.772589
28	3.178054
29	2.772589
30	3.135494
31	3.178054
32	3.135494
33	2.890372
34	2.772589
35	2.890372
36	2.772589
37	3.091042
38	3.091042
39	2.944439
40	3.135494
41	2.944439
42	3.091042
43	2.772589
44	2.833213
45	3.091042
46	2.944439
47	2.772589
48	2.772589
49	2.944439

Name: age, dtype: float64

```
In [ ]: # show age distribution after transformation in boxplot
sns.boxplot(data['age'])
plt.show()
```

